

Predict Album Genre by Cover Image Attributes

John Dawes

November 1, 2019

Business Case

In the music industry, it is not unusual for labels to expect that a newly signed artist/band to have their all their branding ready to go for release at contract signing. This includes the music being mixed and mastered, along with cover art for their first release on that label. It would be useful for the artist to have a tool that would give them an idea of how much “like” their album cover communicates the genre of their music before they shop labels. For labels, to release their artists’ sophomore (and later) albums, it would be a handy tool to input the album cover to get an idea of how much it represents the genre by album cover attributes.

Data Case

Some useful attributes of albums to predict genre would be the **title**, how many **faces** are on the cover, the most prevalent **colors**, identification of various **objects** and how abstract the cover image is - i.e. is it an **illustration** or a **photograph**.

The data used in this analysis comes from several sources:

- *Artist Names* come from my personal music collection and is based on the folder names in a Music folder.
- From *Artist Names*, a search was performed via the Discogs.com search API: <https://api.discogs.com/database/search?artist=>. This yielded the first 50 albums in an artist discography, that each contained: **year**, **title**, **genre**, **format**, and the all important **cover_image** url.
- Once all the album information from my artist collection was scraped, three Google Vision APIs to find how many faces, what color content, if there were any musical objects, and whether or not the cover is an abstract illustration or photograph.

The following three sections describe how these image attributes were aggregated to create the album data from the Google Vision API image search results:

Faces

The **faceDetection** feature of Google Vision returns an object for each face found in an image, where those faces are (via coordinates for the eyebrows, eyes, and nose), and the probability of the emotion detected. Many album covers contain portraits of either a solo artist, the entire band, or a human model. While being able to differentiate what kind of portraits/human subjects are in the photo, the album data only includes the number of faces detected in the image, for the sake of simplifying the use of **faces** as a predictor in **albumData**.

Colors

Based on the **coloraze** npm module to convert Hexidecimal values to the closest color in it’s palette of 2196 colors. While this library is still a great reduction of the 16 Million colors that can be represented in Hex, I reduced the palette to *16 web-safe colors* to distill the essence of the album image colors into more basic color schemes and to simplify the color predictors in **albumData**.

<https://www.rapidtables.com/web/color/color-wheel.html>





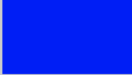



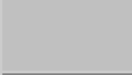







Color	Color name	Hex	(R,G,B)	(H,S,L)
	Black	#000000	(0,0,0)	(0°,0%,0%)
	White	#FFFFFF	(255,255,255)	(0°,0%,100%)
	Red	#FF0000	(255,0,0)	(0°,100%,50%)
	Lime	#00FF00	(0,255,0)	(120°,100%,50%)
	Blue	#0000FF	(0,0,255)	(240°,100%,50%)
	Yellow	#FFFF00	(255,255,0)	(60°,100%,50%)
	Cyan	#00FFFF	(0,255,255)	(180°,100%,50%)
	Magenta	#FF00FF	(255,0,255)	(300°,100%,50%)
	Silver	#C0C0C0	(192,192,192)	(0°,0%,75%)
	Gray	#808080	(128,128,128)	(0°,0%,50%)
	Maroon	#800000	(128,0,0)	(0°,100%,25%)
	Olive	#808000	(128,128,0)	(60°,100%,25%)
	Green	#008000	(0,128,0)	(120°,100%,25%)
	Purple	#800080	(128,0,128)	(300°,100%,25%)
	Teal	#008080	(0,128,128)	(180°,100%,25%)
	Navy	#000080	(0,0,128)	(240°,100%,25%)

Figure 1: 16-Color Palette

Captions

EDA

```
albumCovers <- read_csv('data/album-cover-data.csv')

glimpse(albumCovers)

## Observations: 865
## Variables: 11
## $ artist      <chr> "Bauhaus", "Bonobo", "Delorean", "Haxan Cloak,...
## $ year        <dbl> 2005, 2013, 2010, 2013, 1978, 1997, 2017, 2009...
## $ title       <chr> "New York City, NY 11.12.05", "The North Borde...
## $ genre       <chr> "Electronic", "Electronic", "Rock", "Electroni...
## $ faces       <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0...
## $ primary_color <chr> "White", "Gray", "White", "Black", "Black", "0...
## $ primary_score <dbl> 0.79221308, 0.52558565, 0.53635633, 0.40122879...
## $ secondary_color <chr> "Silver", "Teal", "Silver", "Gray", "White", "...
## $ secondary_score <dbl> 0.08269589, 0.24687445, 0.27873254, 0.29843235...
## $ tertiary_color <chr> "Black", "Gray", "Silver", "Black", "Silver", ...
## $ tertiary_score <dbl> 0.05008276, 0.13973045, 0.13341449, 0.29201144...
```

Machine Learning

Supervised learning is an excellent candidate for training a machine on album attributes, as we already know the `genre`

Analysis